

NEWSFEED SYSTEM

| M | T | W | T | F | S | S |
|-----------|-------|---|---|---|---|---|
| Page No.: | YOUVA | | | | | |
| Date: | | | | | | |

STEP 1: CLARIFICATION OF REQUIREMENTS, ESTABLISH DESIGN SCOPE.

Important Features :-

User: Publish post

See his/her friends posts.

Follow/Unfollow

Like & Comment.

Clarification:-

Posts sorted by

reverse chronological

order? Yes.

Post has media? Yes.

* Functional Requirements:

- Social Media Post - text, video, image
- Follow/ Unfollow.
- Newsfeed.
- Like & Comment
- Notifications

* Non-functional Requirements:

- Availability 99.999%
- Eventual Consistency
- Latency
- Scalability 500M DAU & 2B MAU
- Extensibility
- Reliability
- Fault Tolerance

| Y | T | W | F | S |
|-----|---|---|---|-------|
| Mon | | | | Thurs |
| Tue | | | | Fri |
| Wed | | | | Sat |

STEPS

CAPACITY ESTIMATION

- DAU & MAU
- Throughput
- Storage
- Memory
- Network

fun no. of servers
cost management
divide logical specifications
of all hardware.

→ No. of servers

→ DAU & MAU

1B user | 500M DAU
2B MAU.

→ Throughput

① write throughput

→ comments / likes

→ create post ↑

→ follow/unfollow

(Each user has 300 friends
& follows 250 pages
on average.)

comments/likes activities : each user
performs 3 activities/day

create posts : 10% of DAU post in
a day.

follow/unfollow : 1 user follows 1 user
per week.

② read throughput

a user reads

→ 100 post in a day.

500M users x 100

= 50B user read /day.

∴ 50B read requests/day.

→ storage

- posts
- follows
- activities (likes / comments)
- user metadata

① user metadata

50 KB/user.

$$1B \times 50KB = 50TB \times 8 \times \text{Month}$$

② posts

50M posts/day

20% text

text - 100 KB

20% video

video - 20 MB

60% images

image - 0.5 MB

$$\begin{aligned} \text{Text storage} &= 50M \times 0.2 \times 100KB / \text{day} \\ &= 1TB / \text{day}. \end{aligned}$$

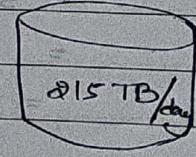
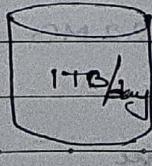
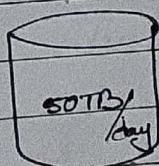
$$\begin{aligned} \text{Image storage} &= 50M \times 0.6 \times 0.5MB / \text{day} \\ &= 15TB / \text{day} \end{aligned}$$

$$\begin{aligned} \text{Video storage} &= 50M \times 0.2 \times 20MB / \text{day} \\ &= 200TB / \text{day}. \end{aligned}$$

User metadata

Text posts

media content



(3) follows.

1 follow - 10B, 1 user follows 1 user per week.
 $\frac{(500M \times 10B)}{7}$ per day.

$\left(\frac{500M \times 10B}{7} \times 365 \times 10 \right)$ Bytes → for 10 years.

(4) activities.

1 activity - 216B, 3 activities per user per day.

$(500M \times 3 \times 216B \times 365 \times 10)$ Bytes → for 10 years.

→ Memory (cache)

cache = 1% of storage

storage = $0.01 \times 216 \text{ TB/day}$

= 2.16 TB/day .

→ Network or Bandwidth Estimation

216TB/day - storage

50B read requests/day.

* $\frac{216\text{TB}}{86400} / \text{sec.} \rightarrow \text{ingress}$

= 2.5 GB/sec.

* $50B * \text{avg. size of post}$

= $50B * (0.2 * 100\text{KB} + 0.2 * 20\text{MB} + 0.6 * 0.5\text{MB})$

= $50B * 4.32\text{MB}$

= 216 PB/day.

= $\frac{216 \text{ PB}}{86400} / \text{sec.} = 2.5 \text{ TB/sec.}$

→ No. of servers.

$$\begin{aligned}
 \text{Peak load} &= 500 \text{M/sec.} \\
 \text{1 server capacity} &= 64000 \text{ RPS.} \\
 \text{No. of servers} &= \frac{500}{64000} \\
 &= 7812.5 \\
 &\approx 8K \text{ servers.}
 \end{aligned}$$

STEP 3: API & HIGH LEVEL DESIGN.

- * Database
- * Cache
- * Blob storage
- * CDN
- * Load Balancers

POST /v1/posts {
 user_id :
 media_url :
 text :
 hashtags :
 }
 }

POST /v1/comments {
 user_id :
 post_id :
 comment :
 }
 }

POST /v1/likes {
 user_id :
 post_id :
 }
 }

To update like/comment:

PATCH: /v1/posts/{postId} }
body: like or comment data. } extendable

(8)

POST: /v1/comments and /v1/likes } scalable.
body: { }

POST: /v1/follow.

body: { }

example: { "followerId":
"followerId":
"followeeId":
"followeeId": }

Read newsfeed:-

GET: /v1/Feed?access_token={}

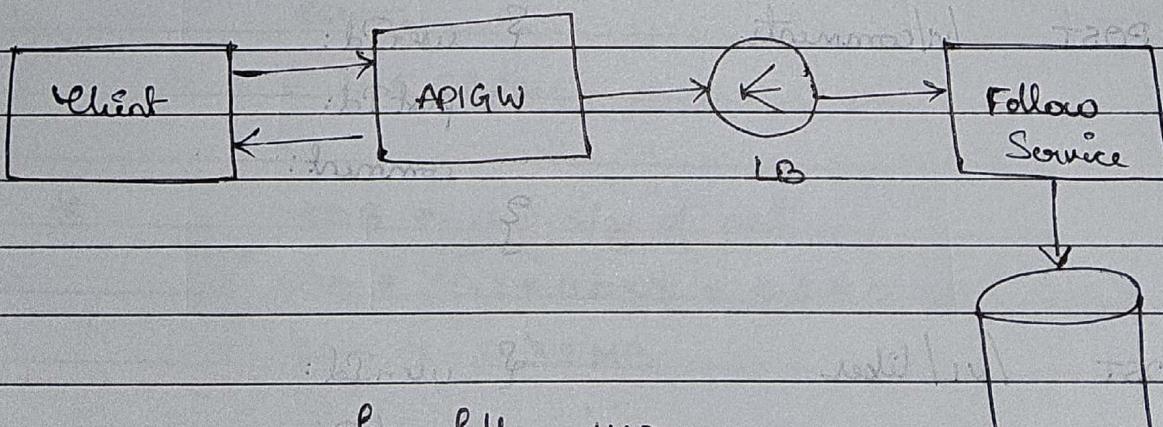
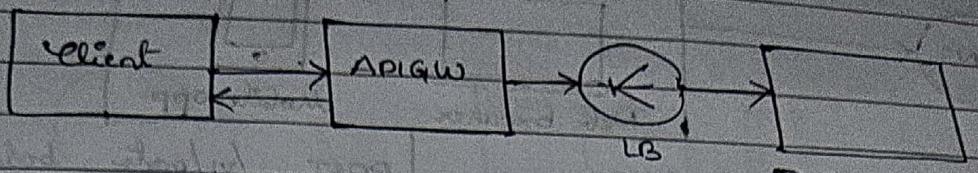
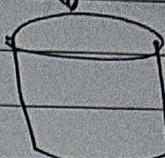


fig. - follow HLD

-Follow DB.

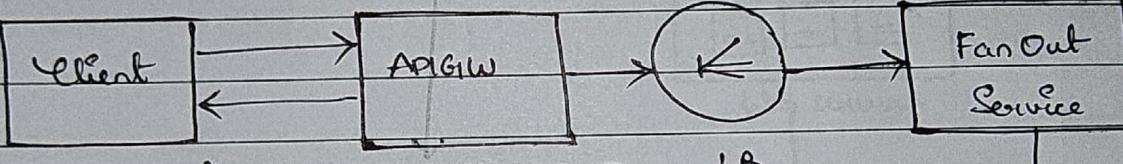


Post Service

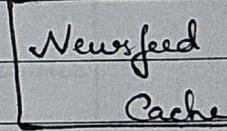


Post DB

Fig, Create posts



Fan Out Service



Newsfeed Cache

Fig, Feed publishing.

create job
service 2

transit
service

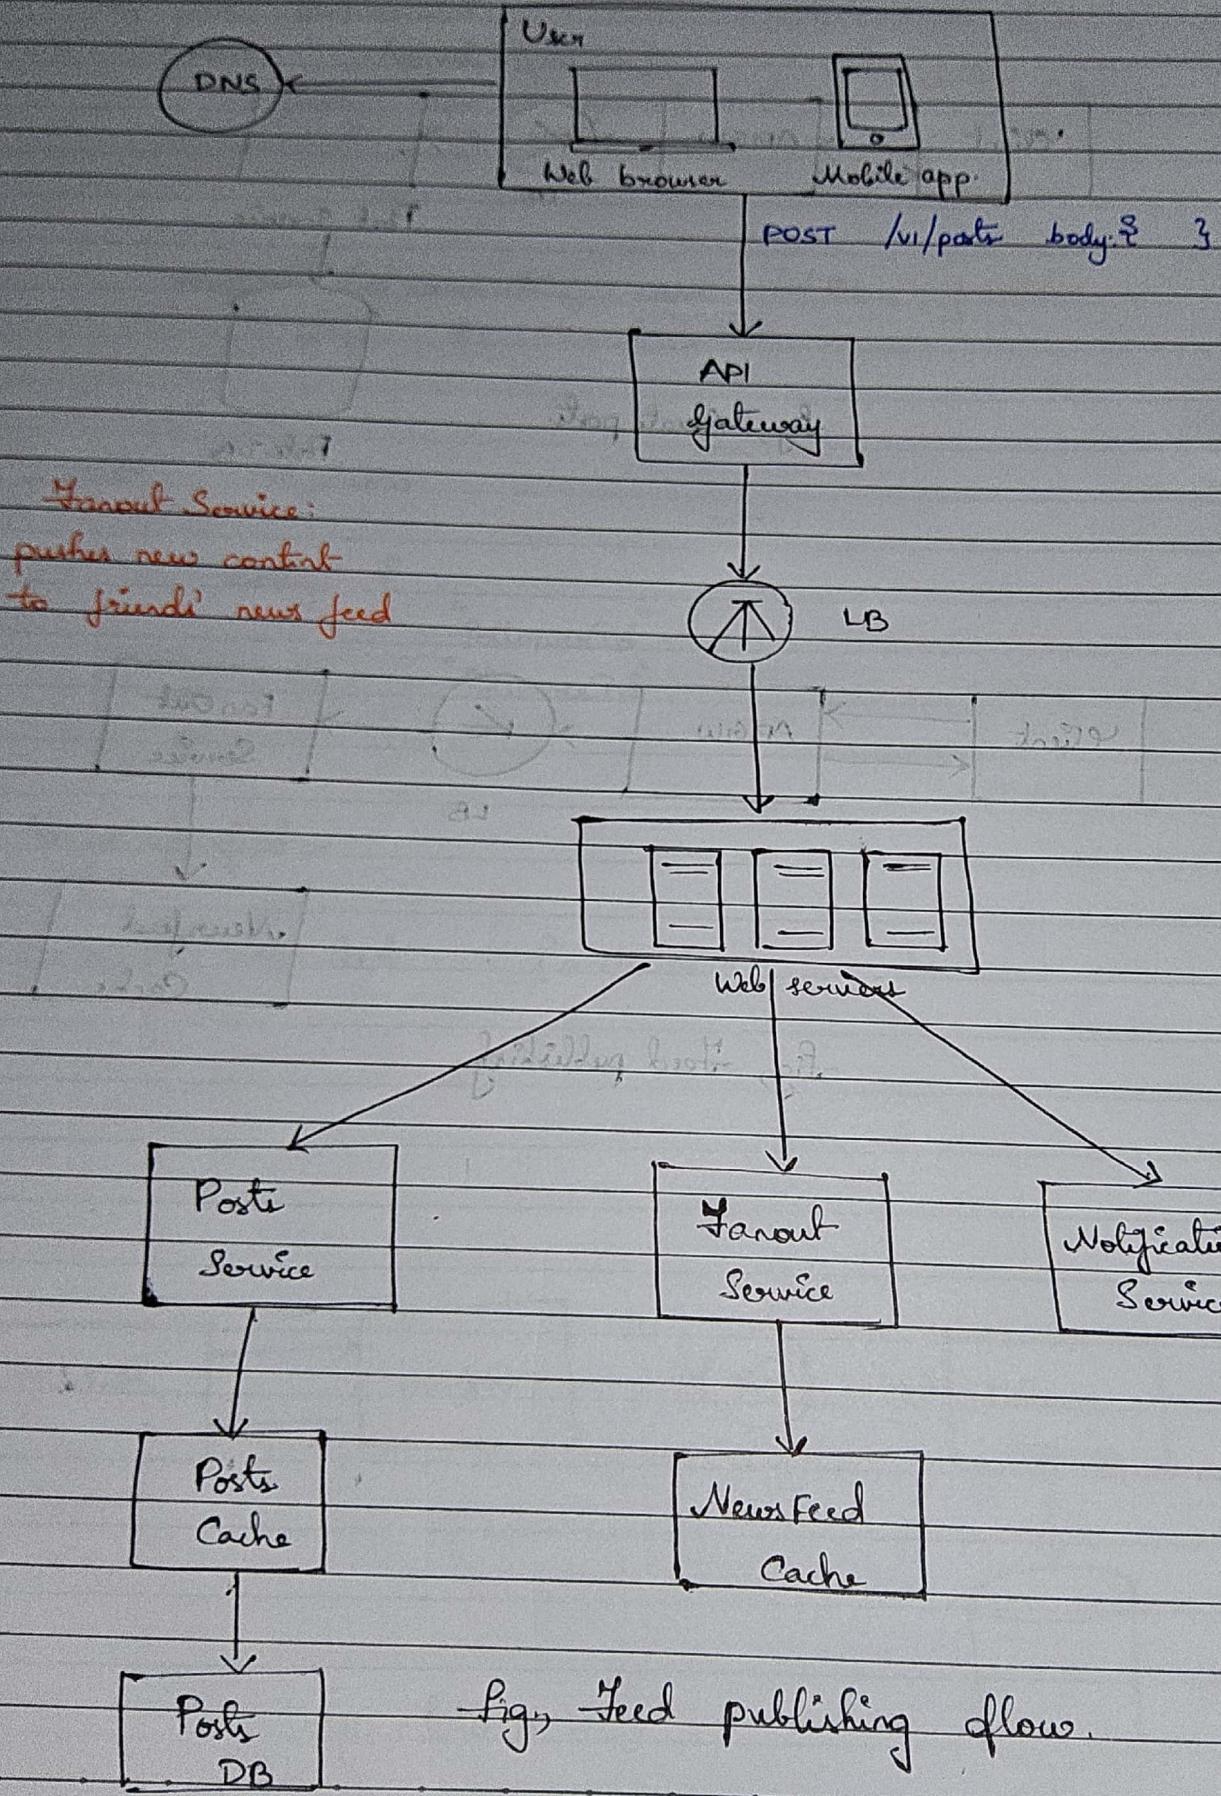
host
service

frequency
info

start
info

snappy publishing info

etc



Newsfeed Building :-

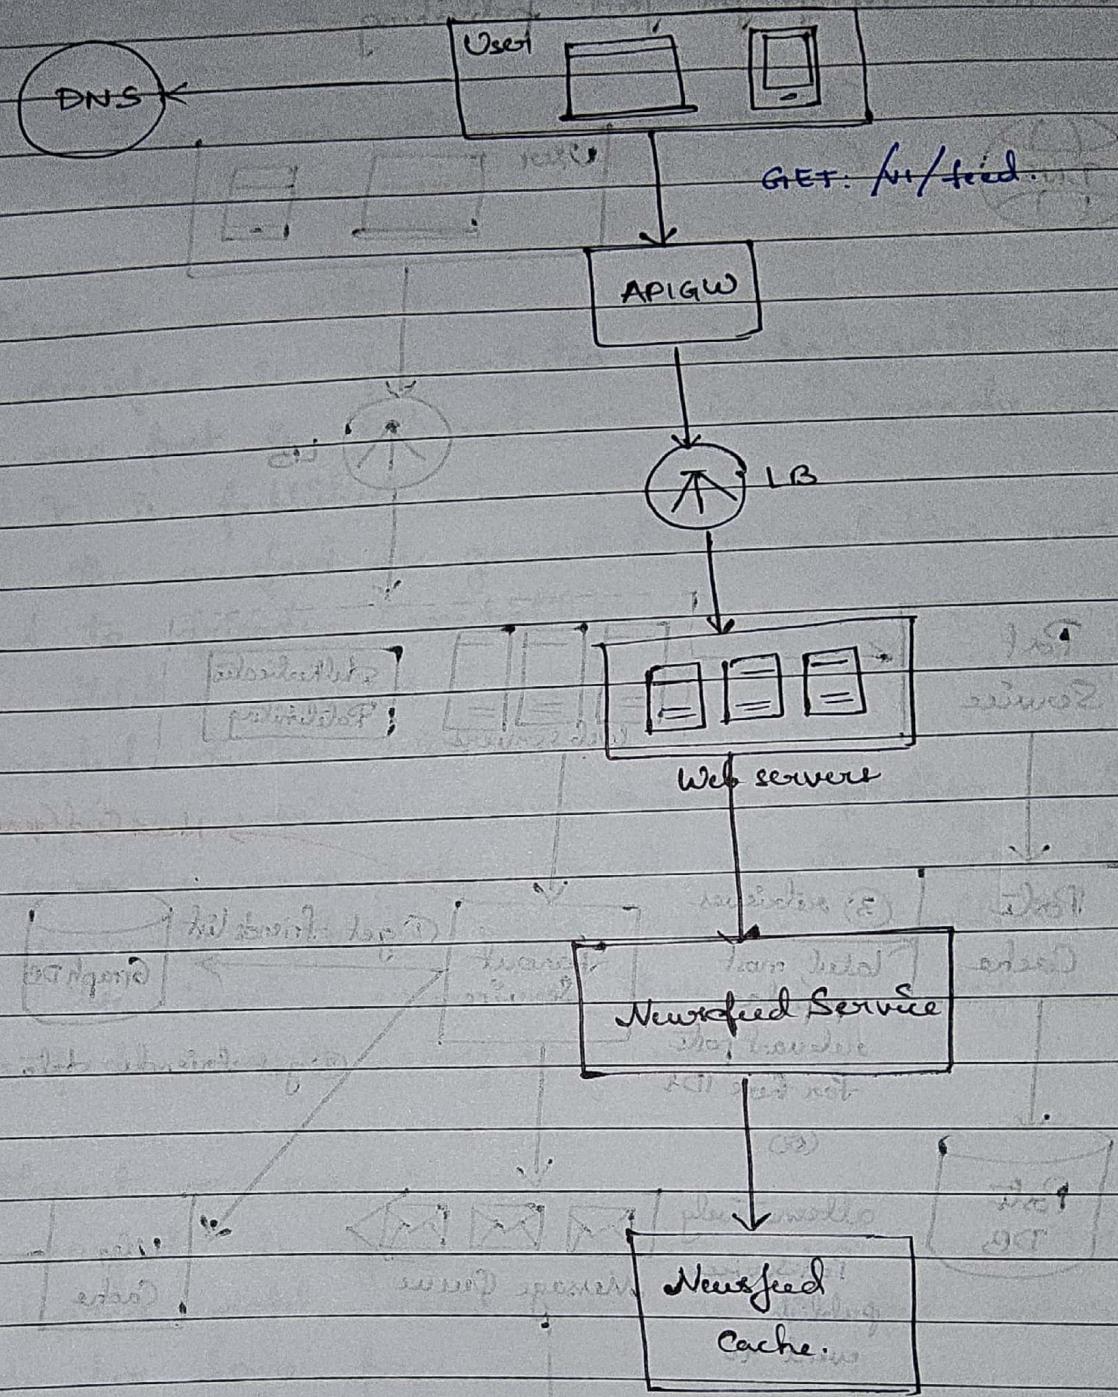
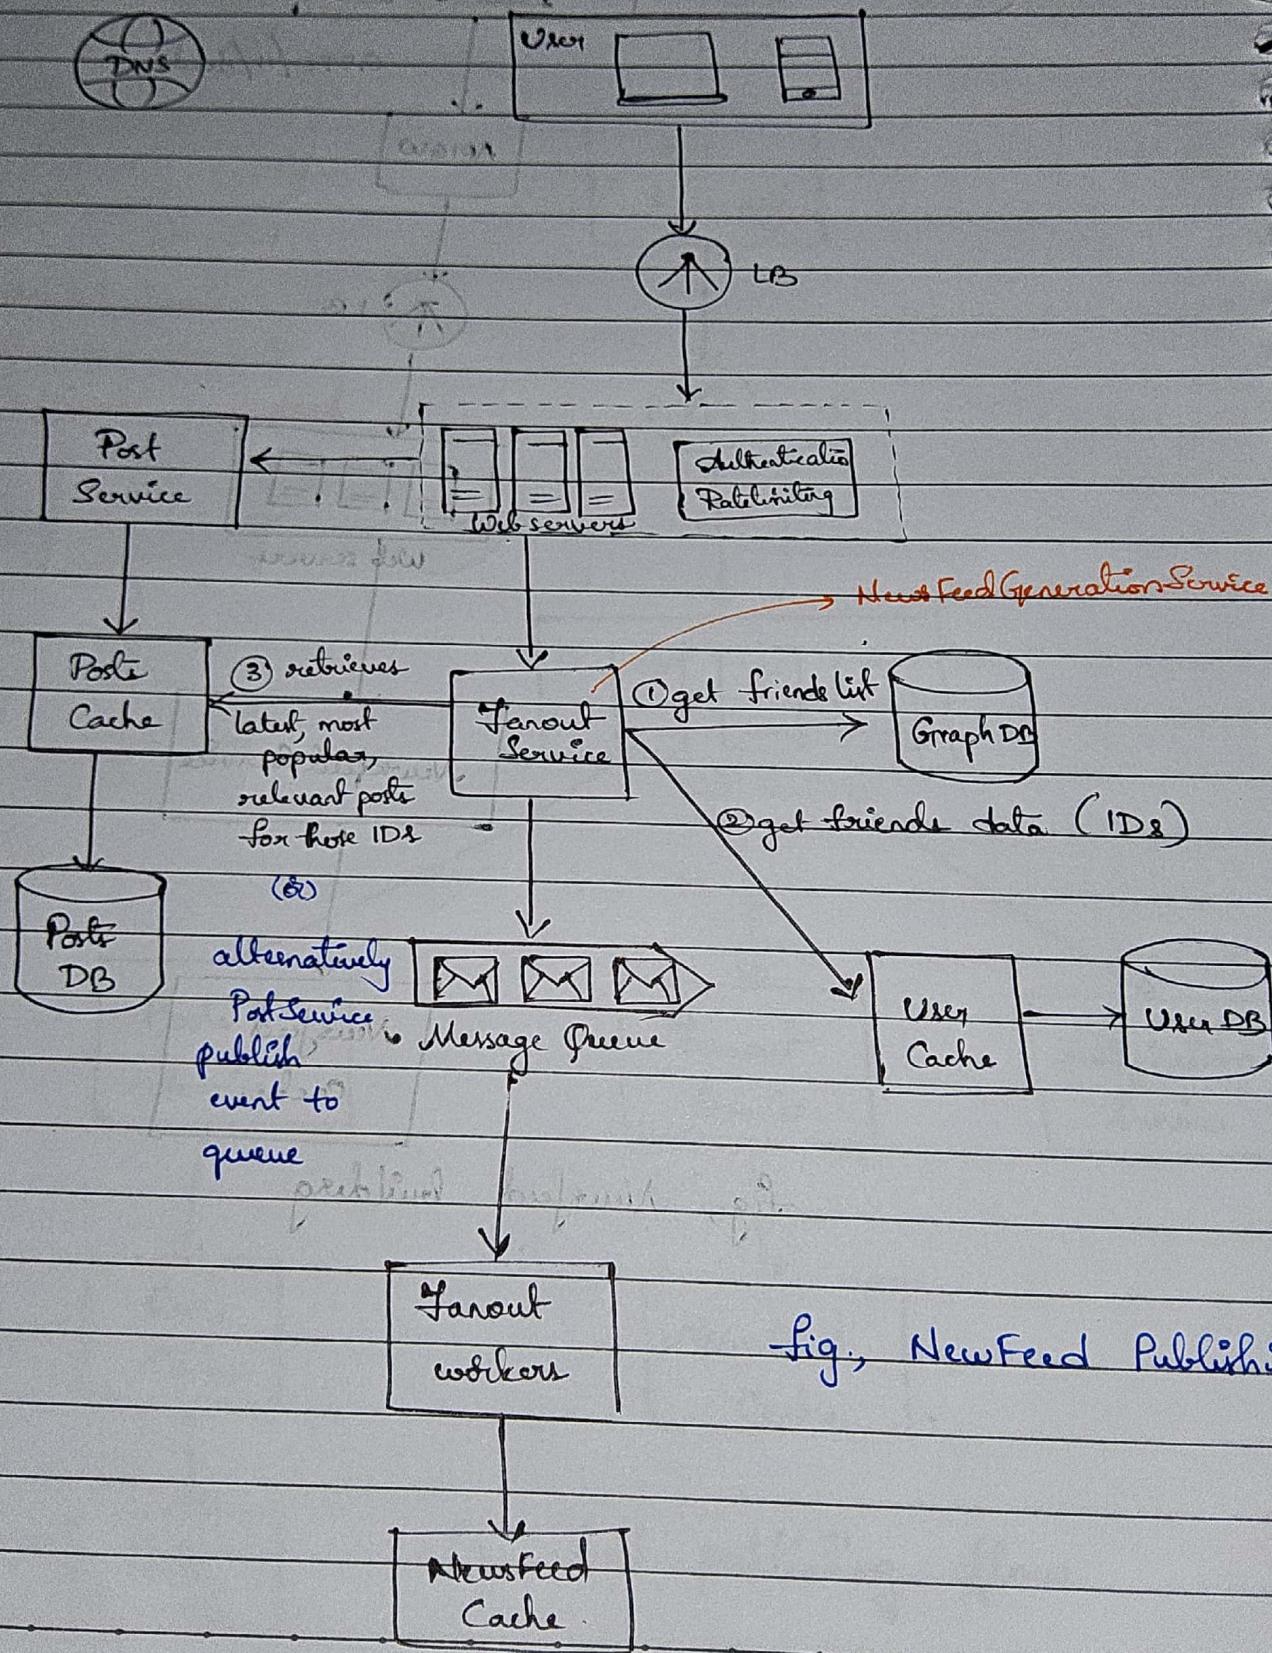


Fig., Newsfeed building

STEP 4 :- DESIGN DEEP DIVE

Detailed design for feed publishing :-



Fanout :-

Fanout is the process of delivering a post to all friends.

- Fanout on write (push model)
- Fanout on read (pull model)

* Fanout on write :-

newsfeed is precomputed during write time.
A new post is delivered to friends' cache immediately after it is published.

Pros: The newsfeed is generated in real-time and can be pushed to friends immediately.

Cons: Fetching newsfeed is fast because the newsfeed is precomputed during write time.

Cons: If a user has many friends, fetching friends list and generating news feeds for all of them are slow and time consuming. It is called hot-key problem.

For inactive users or those rarely log in, pre-computing news feeds waste computing resources.

* Fanout on read :-

newsfeed is generated during read time. This is an on-demand model. Recent posts are pulled when a user loads her home page.

Pros: For inactive users or those who rarely log in, fanout on read works better because it will not waste computing resources on them.

Data is not pushed to friends so there is no hot-key problem.

Cons: Fetching the news feed is slow

We adopt a hybrid approach.

When a celebrity user posts: Fanout-on-Read

→ Post is stored in Posts DB.

→ Post is added to 'Global Cache' & 'CDN', so it can be fetched on demand.

→ Follower opens their feed - get normal posts from their newsfeed cache.

→ check if they follow any celebrity - fetch their posts from 'Global Cache'.

→ Merge both sets of posts dynamically.

When a normal user posts: Fanout-on-Write

→ Post is stored in Posts DB.

→ Fetch the list of followers from Followers DB.

→ push post_id, user_id into newsfeedcache.

↓
follower_id fetched from DB.
get from queue where post service publishes event
& get from post cache.

→ Follower opens their feed - directly fetches posts from newsfeed cache

Another option is while publishing we only push to online followers after a "normal" user publishes a post.

-Flow of Newsfeed Publishing (see in fig.) :-

User posts \rightarrow /v1/post (POST)

Request flows through load balancer, servers & goes to post service.

Post Service saves the post in Post DB and updates the Post Cache.

post id is generated.

post-id, author_id & data flows to three Fanout Services (or Newsfeed generator service)

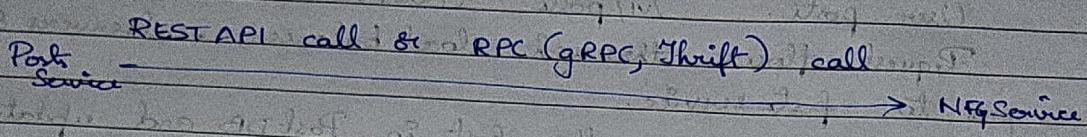
Newsfeed generator (or Fanout) service works as follows:

1. Fetch friend IDs from the graph database (Follows DB).
2. Get friends info from the user cache. The system then filters out friends based on user settings.
For example, if you mute someone, his posts will not show up on your news feed even though you are still friends.
3. Send friends list and new post ID to the message queue.
4. Fanout workers fetch data from the message queue and store news feed data in the news feed cache.
5. Append each \langle post-id, user-id \rangle to news feed cache

| | |
|---------|---------|
| post_id | user_id |
| post_id | user_id |
| post_id | user_id |

Methods for sending post id to Newsfeed generator:-

① Synchronous API call (Direct approach)



After post service stores the post in Post DB, returns post id

Post service calls POST /newsfeed/fanout

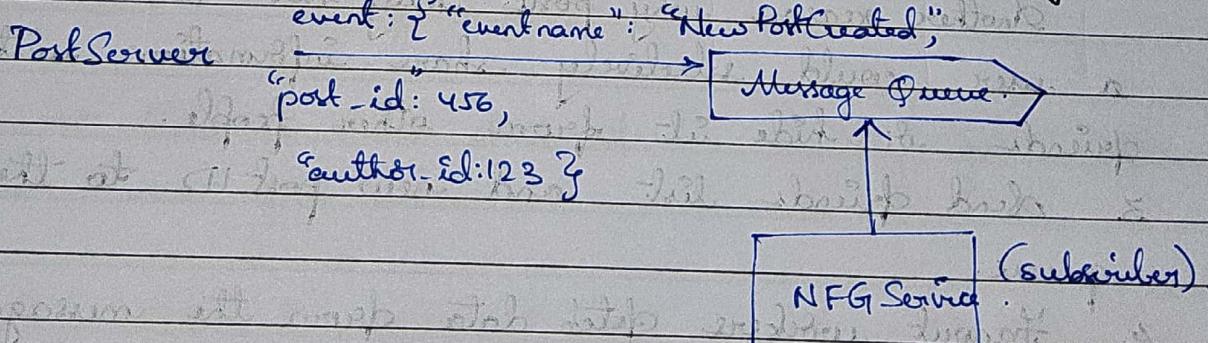
body: { "post_id": "456",
"author_id": "123" }

NFG service fetches followers by author_id & pushes {post_id, follower_id} to queue.

Cons:-

- ↳ Increases latency
- ↳ If Newsfeed generator fails, retries are needed.

② Asynchronous Event-Driven (Better for Scalability).



Pros:- Decoupled & Scalable.
Can handle failures via retries.

③ Database Polling (from Posts Cache).

Create post:

User → POST /feed/posts : content
Post Service saves post in Post DB & generates post_id

Post Cache updated:

Post Service also updates Post Cache

e.g., SET post:123 (latest_post_id = 456) EX 3600

NFG service retrieves post_id from post cache:

e.g., GET post:123

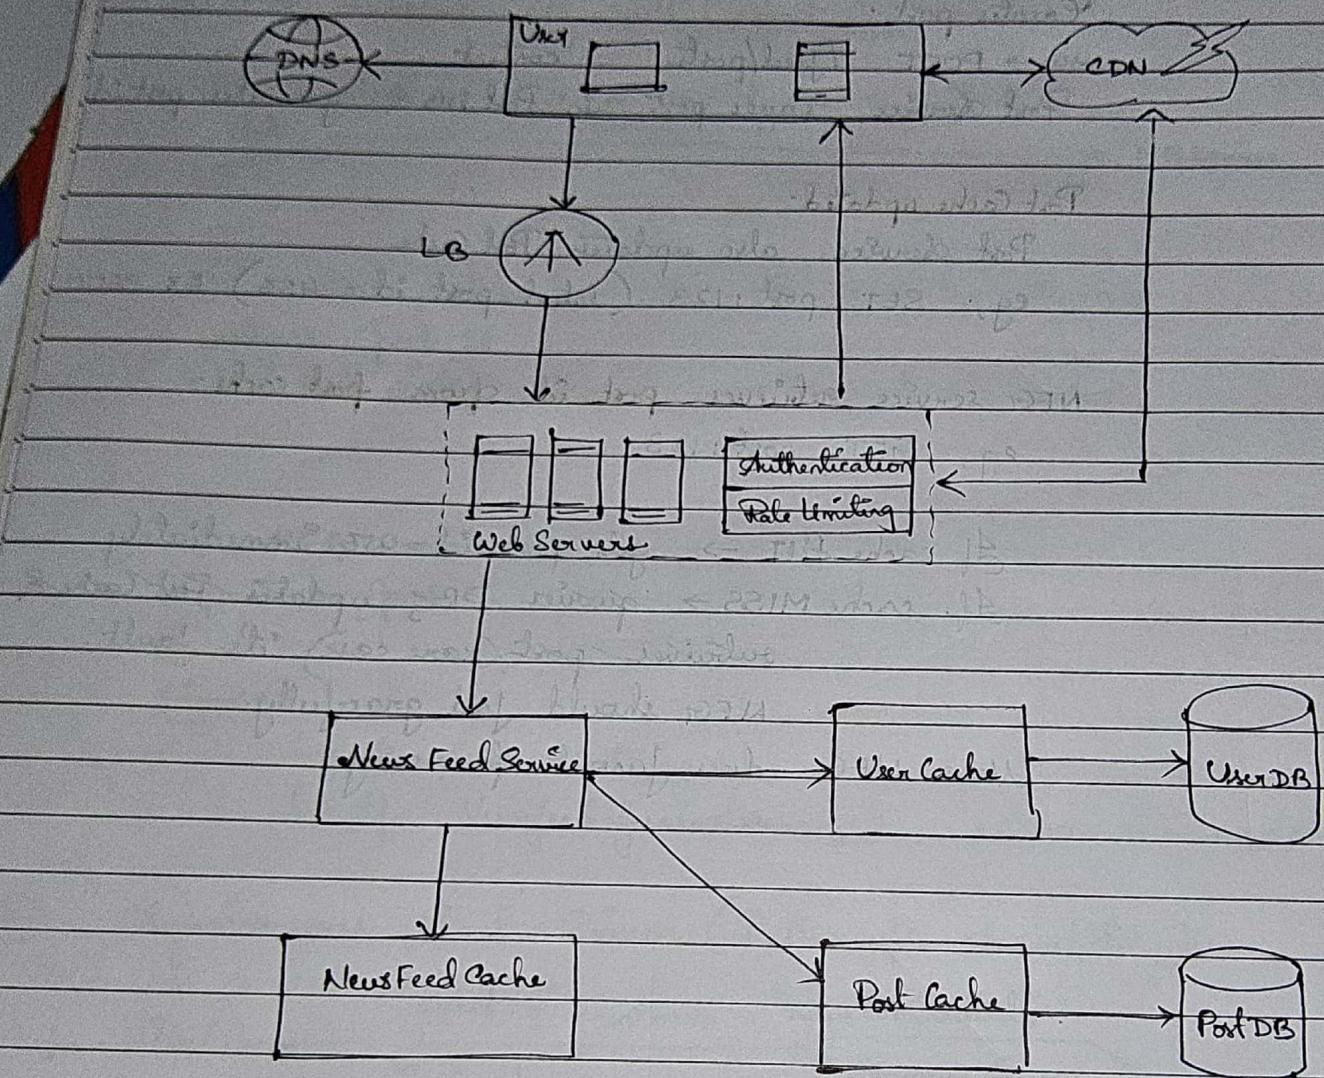
If cache HIT → gets post_id = 456 immediately

If cache MISS → queries DB, updates Post Cache & retrieves post (rare case) it's fault.

NFG should fail gracefully.

Now NFG does Janout processing.

Detailed Design for Newsfeed retrieval:-



1. A user sends a request to retrieve her news feed.
2. Load Balancer redistributes requests to web servers.
3. Web servers call the newsfeed service to fetch news feeds.
4. News feed service gets a list of post IDs from newsfeed cache.
5. News feed service fetches posts related content & user content from User Cache, Post Cache to construct a fully hydrated newsfeed.

6. The fully hydrated newsfeed is returned in JSON format back to the client after rendering.

What is the problem with generating a newsfeed upon a user request? (also called live update)

Suppose if our service receives billions of requests at once and the system starts generating live newsfeeds.

This situation would put an enormous number of users on hold and could crash servers. To avoid this, we can assign dedicated servers that will continuously rank and generate newsfeed and store them in the newsfeed database and memory.

Therefore, upon a request for new posts from a user, our system will provide the pre-generated feeds.

* The newsfeed ranking service:-

After we see relevant and important posts on the top of our newsfeed whenever we log in to our social media accounts.

This ranking involves multiple advanced ranking and recommendation algorithms.

"Newsfeed ranking service" consists of these algorithms working on various features, such as, a user's past history, likes, dislikes, comments, clicks, and many more.

These algorithm also perform the following functions:

- ① Select "candidate" posts to show in a newsfeed.
- ② Eliminate posts including misinformation and clickbait from the candidate post.
- ③ Create a list of friends a user frequently interacts with.
- ④ Choose topics on which a user spent more time.

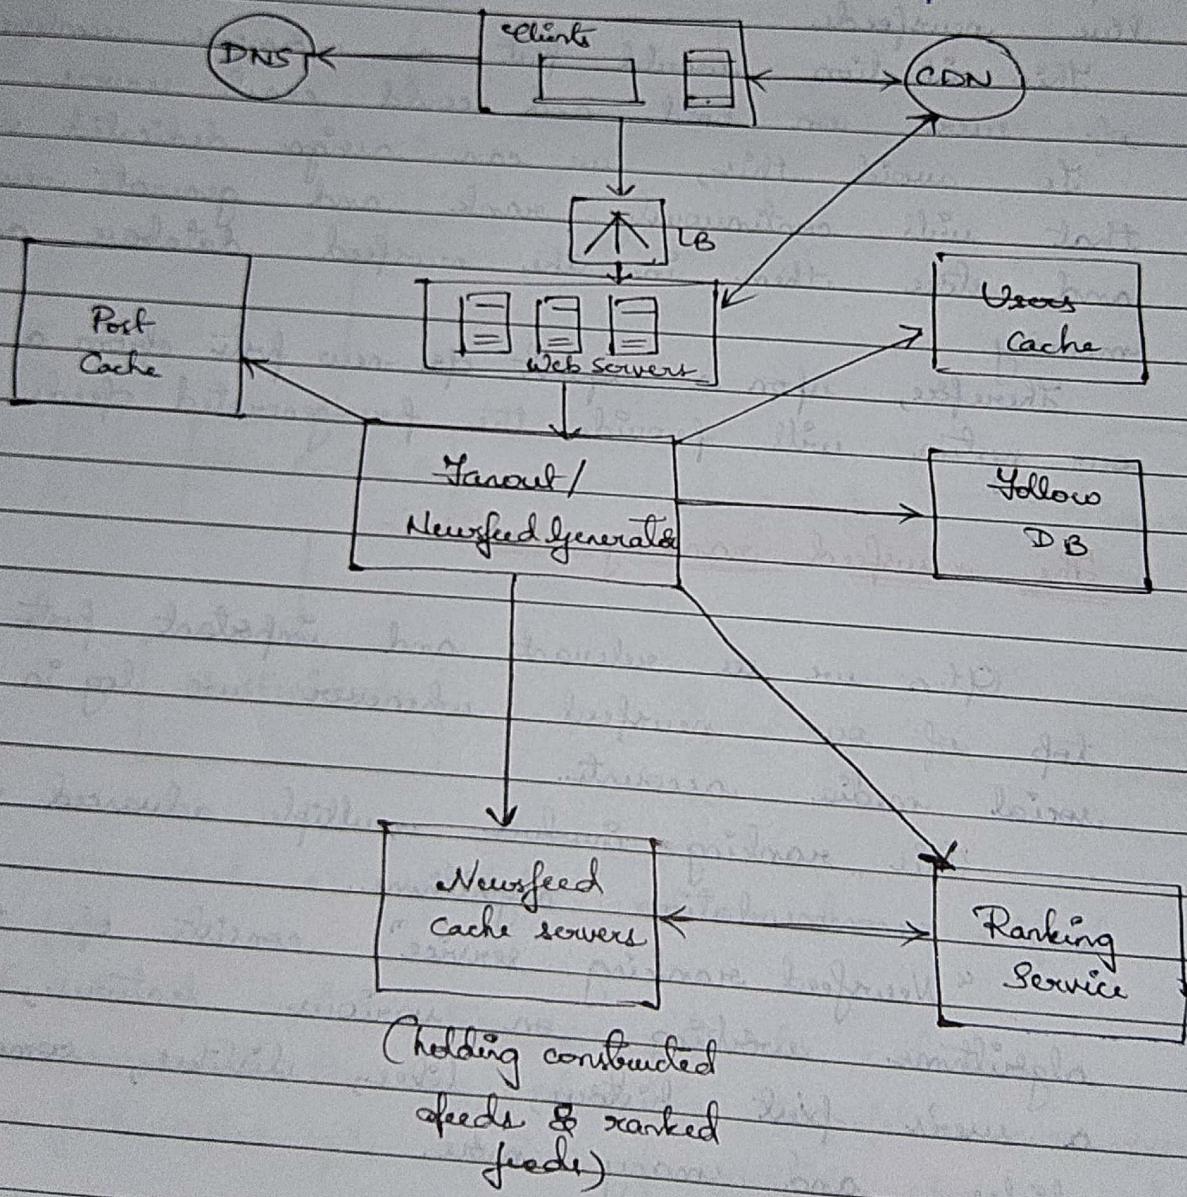


Fig. Feed publishing with Ranking

* Post ranking & Newsfeed construction :-

Assume there are 10 posts in the database - published by 5 users whom Bob follows.

We aim to rank only 4 posts out of 10.

We follow below steps for ranking and creating a newsfeed for Bob:

- ① Various features such as like, comments, shares, category, duration, etc and so on, are extracted from each post.
- ② Based on Bob's previous history, stored in userDB the relevance is calculated for each post via different ranking and machine learning algorithms.
- ③ A relevance score (1 to 5) is assigned to each post.
- ④ Top 4 out of 10 are selected based on the assigned scores.
- ⑤ Top 4 posts are combined and presented on Bob's timeline in decreasing order of the score assigned.

* Uploading media :-

It is not efficient to send image/video in a request.

First client uploads an image to application, application stores it in an object storage and returns file location.

Client uses that location and sends in his request.

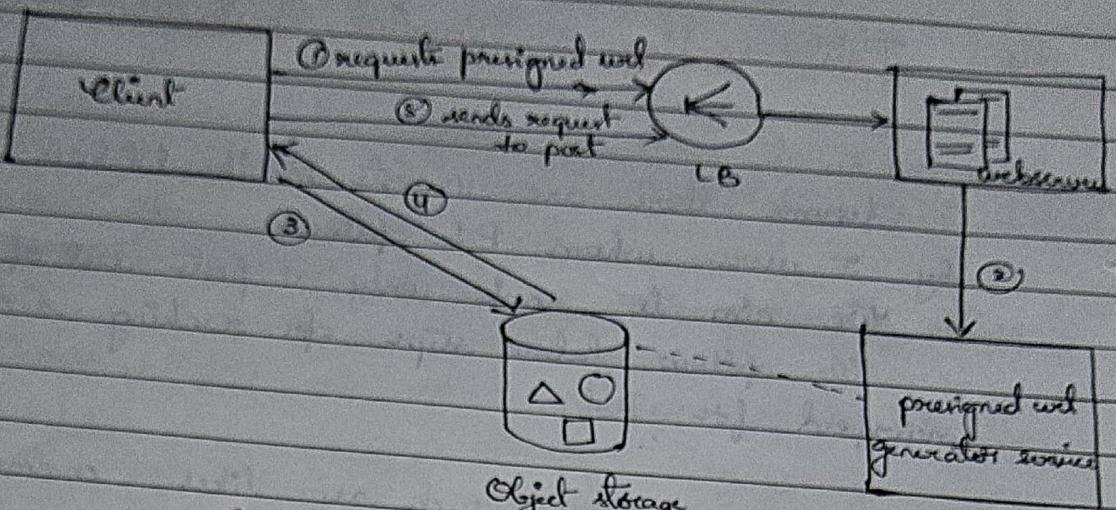
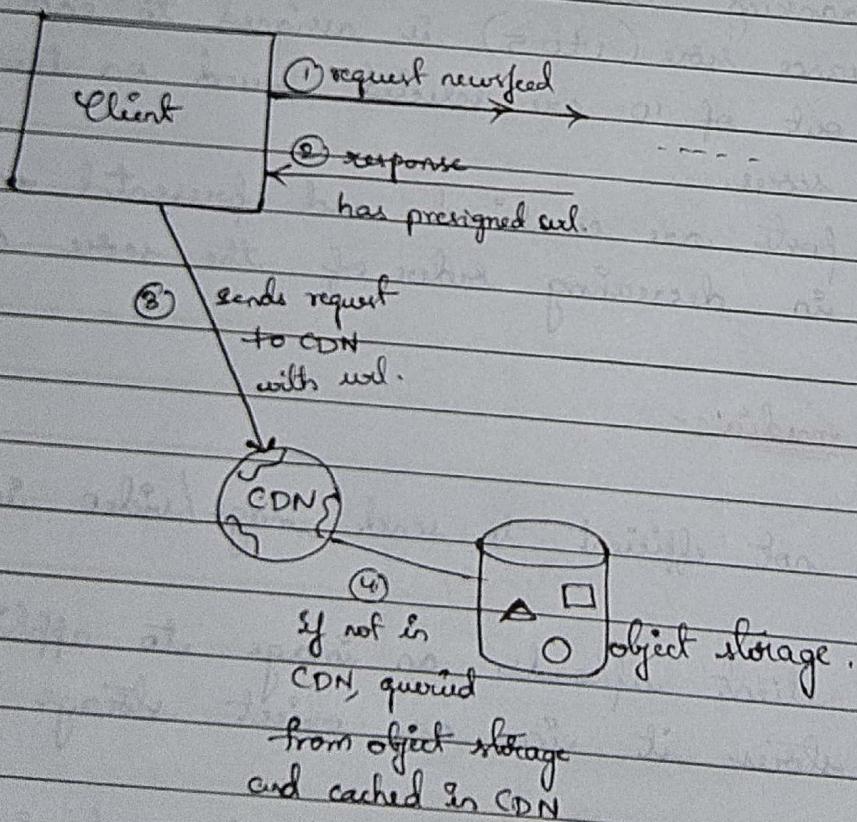


Fig., media upload.

Pre-signed url is a special url that allows client to upload file.

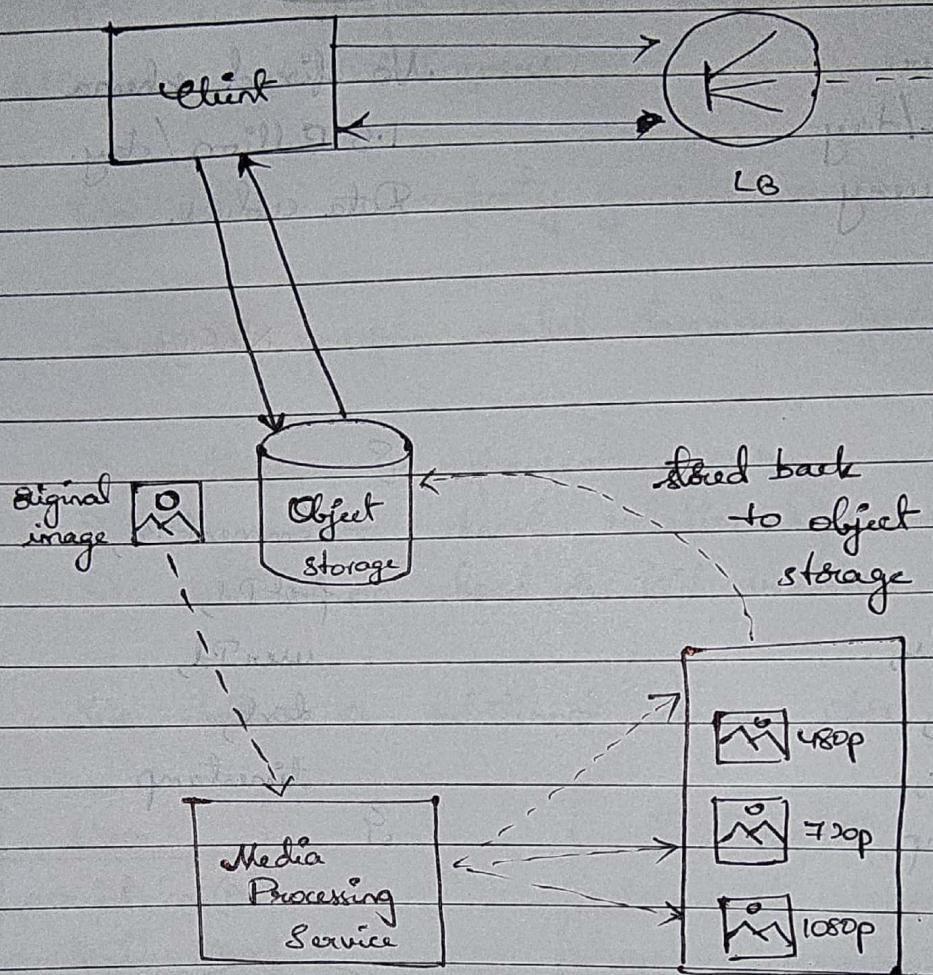


from object storage and cached in CDN

* Media Processing :-

Different formats are required mp4, mov, etc

Different resolutions are required 480p, 1080p etc.



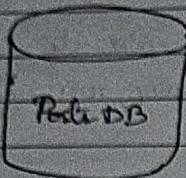
* Database Selection :-

- ① Fast Data access
- ② Scale is too large.
- ③ Data is in fixed structure
- ④ Complex Queries
- ⑤ Data changes frequently / evolves over time

SQL

NoSQL





Text
Metadata
media urls etc

No structure
50M posts/day
Simple query

Nosql



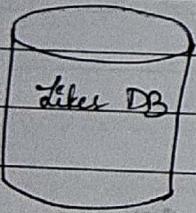
Comments
DB

No fixed schema
1.5 Billion / day.
Data evolves.

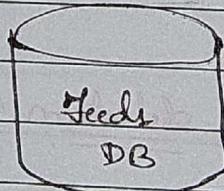
Nosql

postId,
userId,
mediaUrl,
hashtag,
postText,
timestamp

commentId,
postId,
userId,
text,
timestamp



Nosql



Feeds
DB

FeedItems:
[]

No structure

50M posts/day

Simple query pattern

postId,
userId

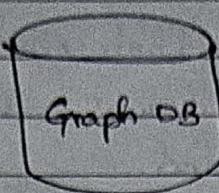
No fixed schema

1.5 Billion / day.

Data evolves.

likeId,
userId,
postId,
timestamp

Follow DB



We use graph DB to store relationships between followers and followees.

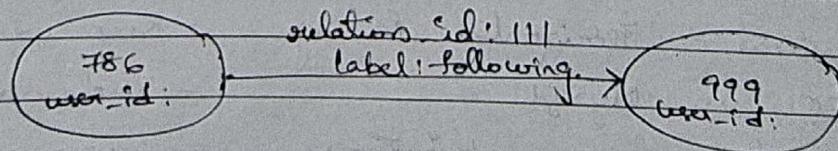
We follow property graph model

Property graph model describes graph by its

- 1) vertices — represent users.
- 2) edges — denote relationships between them.
- 3) indices on head & tail vertices

We follow a relational schema for the graph store

| User | Relationship |
|---------------------------------------|--|
| user_id: int (PK) properties: JSON | relation_id: int (PK) from: int REFERENCES User to: int REFERENCES User (User-ID) label: varchar (32) properties: JSON |



relation_id relation_from relation_to label properties
111 786 999 following JSON-formatted data

STEP 4: EVALUATION AND WRAP UP.

- ① Scalability:- load balancers, web servers and other relevant servers are added/removed on demand.
- ② Fault tolerance :- replication of data redundant resources to handle failures monitoring service to enhance reliability by continuously observing system health, detecting issues early, providing insights for optimization and assisting in timely incident response.
- ③ Availability :- highly available by redundant servers & replicating data when a user gets disconnected due to some fault in the server, the session is re-created via a load balancer with a different server. Data is stored on different and redundant database clusters, which provides high availability and durability.
- ④ Low latency :- we can minimize system's latency at various levels by:
 - geographically distributed servers & cache associated with them.
 - Use CDN for frequently accessed news feeds and media content.
- ⑤ we can talk about
 - Vertical & Horizontal scaling
 - SQL vs NoSQL
 - Master-slave replication
 - Read replicas
 - Consistency models.
 - Database sharding
 - Keep web tier stateless
 - Cache data as much as you can
 - Support multiple DCs.
 - Lose couple with msg q's
 - Monitor key metrics - QPS at peak, latency